

# Ordis: Emergent Intelligence in GPU-Accelerated Virtual Ecosystems

---

## Abstract

---

We present Ordis, a GPU-accelerated framework for simulating emergent intelligence in virtual ecosystems. Using a Four-Key Architecture (adaptive meta-rules, composite action compilation, heritable memory, symbolic communication) and Four-Layer Memory System, we demonstrate genuine behavioral emergence in 200-agent simulations across 20 seeds  $\times$  5000 steps. Key findings: (1) **Antifragility**: 76% pass rate under periodic stress vs 6% in stable conditions ( $p < 0.001$ ); (2) **Multi-lineage civilization**: 36 concurrent lineages, 7 generations, 194 emergent strategies; (3) **Welfare Trap**: low-pressure environments cause behavioral homogenization and communication degradation. We discuss implications for artificial life research and potential risks of immersive applications.

**Keywords:** Emergent Intelligence, Antifragility, Multi-Agent Systems, Artificial Life, GPU Acceleration

**System Version:** Ordis V3.5.33 (main results) / V3.5.39 (extended metrics) | **Data:** 20 seeds  $\times$  5000 steps  $\times$  2 conditions

---

## 1. Introduction

---

### 1.1 Background

Emergence—the phenomenon where complex patterns arise from simple local interactions—remains one of the most fascinating yet elusive concepts in complexity science. From flocking birds to market dynamics, emergent behaviors exhibit characteristics that cannot be predicted from individual component properties alone. In computational systems, achieving genuine emergence (as opposed to pre-programmed complexity) requires careful architectural design that permits bottom-up pattern formation while maintaining system stability. The central challenge lies in creating conditions where novelty can arise spontaneously without designer intervention, yet the system remains robust enough to sustain and build upon these innovations.

**Academic Lineage:** Emergent AI is not a novel concept, but an integration and advancement of established research traditions:

Field	Representative Works	Ordis's Inheritance and Extension
Artificial Life (ALife)	Tierra (Ray, 1991), Avida	Inherits evolutionary dynamics, <b>adds GPU acceleration + social interaction</b>
Multi-Agent Systems (MAS)	Game Theory, MARL	Inherits interaction frameworks, <b>adds adaptive meta-rules</b>
Complex Adaptive Systems (CAS)	Santa Fe Institute (Holland, 1992)	Inherits emergence theory, <b>adds quantifiable verification</b>

The artificial life (ALife) field has a rich history of attempting to create emergent complexity. Ray's Tierra (1991) demonstrated self-replicating programs evolving in digital memory, while Avida extended this with more sophisticated mutation and selection mechanisms. Polyworld (Yaeger, 1994) introduced embodied agents with neural networks in a 3D environment, achieving emergent foraging and mating behaviors. More recently, Neural MMO (Suarez et al., 2019) scaled multi-agent learning to thousands of agents, and Lenia (Chan, 2019) showed continuous cellular automata capable of producing life-like patterns. However, these systems typically focus on either evolutionary dynamics OR social interaction, rarely both simultaneously. Ordis integrates four previously isolated mechanisms—adaptive institutions, behavioral compilation, heritable memory, and symbolic communication—into a unified framework, enabling the study of civilization-level emergence at computational scales.

## 1.2 Research Questions

1. Can complex social structures emerge from simple individual rules in finite-scale simulations?
2. How does environmental pressure affect emergent behavior and system stability?
3. What architectural components are necessary for sustained behavioral innovation?

## 1.3 Contributions

1. **Four-Key Architecture** for emergent intelligence (H1-H4)
2. **Antifragility validation:** 76% pass rate under stress vs 6% stable ( $\chi^2 = 15.48$ ,  $p < 0.001$ )
3. **Ablation study:** Social mechanisms (SHARE/STEAL) causally drive Matthew Effect (-12.6% Gini when disabled)
4. **Capacity saturation solution:** Fitness-based culling enables sustained multi-generation evolution (gene\_moran: 0 → 0.17)
5. **Long-term stability:** 7 generations, 5,000 steps without collapse
6. GPU-accelerated simulation achieving 675x speedup

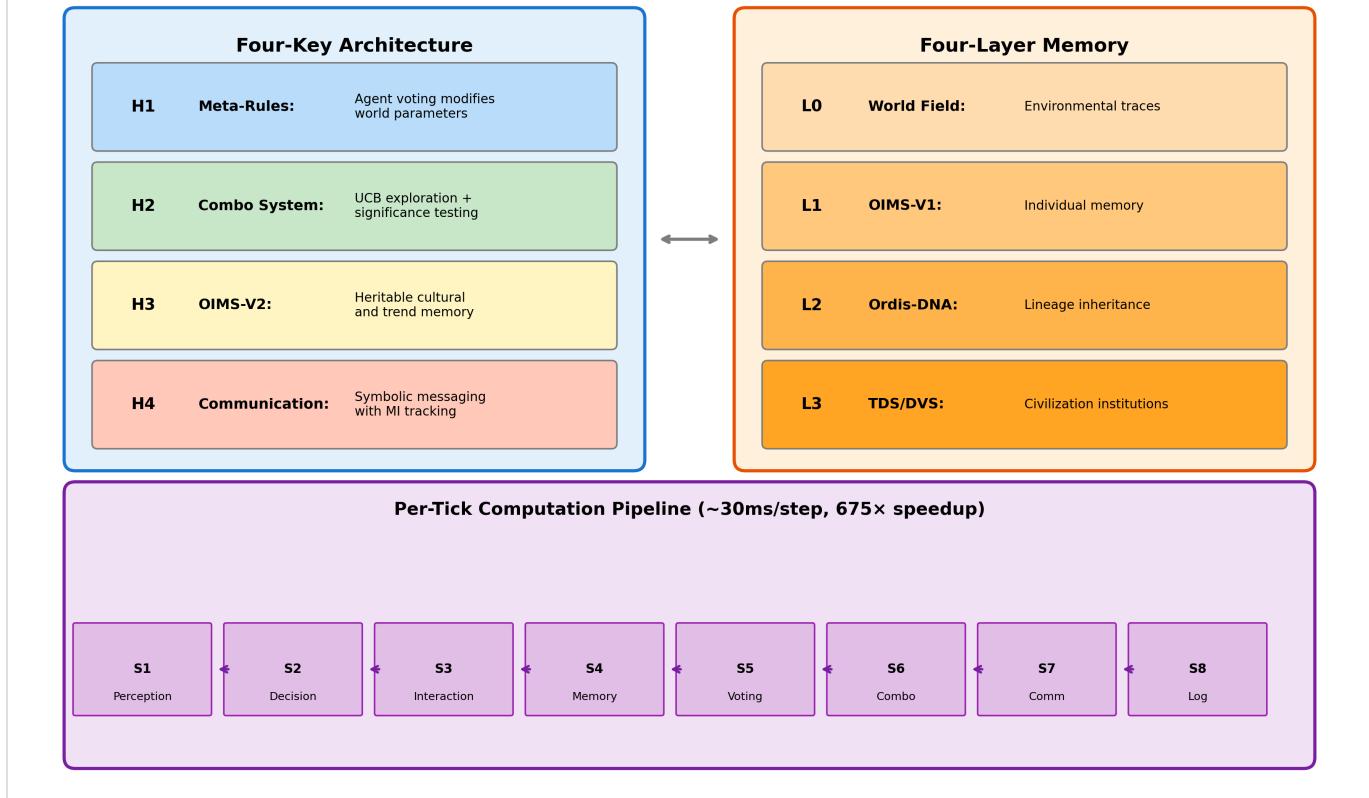
---

## 2. Methods

---

### 2.1 System Architecture

## Ordis System Architecture



### 2.1.0 Per-Tick Computation Pipeline

Each simulation step executes the following 8-stage pipeline:

- Stage 1: Perception Aggregation
  - └ Entities sense local resources, neighbors, field state
- Stage 2: Decision Sampling
  - └ Action selection via UCB exploration + learned combos
- Stage 3: Interaction Settlement
  - └ GATHER/STEAL/SHARE resolution, energy/resource transfer
- Stage 4: Memory Write (L1)
  - └ OIMS-V1 individual memory update
- Stage 5: Meta-Rule Voting (H1)
  - └ Threshold adjustment votes aggregated and applied
- Stage 6: Combo Compilation (H2)
  - └ Pattern detection, synergy testing, acceptance/rejection
- Stage 7: Communication Metrics (H4)
  - └ Message-action mutual information computation
- Stage 8: Snapshot & Logging

```
|   └ Every 200 steps: tick_agg.jsonl; events: events.jsonl |
```

This pipeline is fully tensorized on GPU, achieving ~30ms/step (675x vs CPU baseline).

### 2.1.1 Four-Key Architecture

Key	Component	Function
H1	Meta-Rules	Agent voting modifies world parameters
H2	Combo System	UCB exploration + significance testing for strategy discovery
H3	OIMS-V2	Heritable cultural and trend memory
H4	Communication	Symbolic messaging with mutual information tracking

### 2.1.2 Four-Layer Memory

Layer	Name	Scope
L0	World Field	Environmental traces
L1	OIMS-V1	Individual memory
L2	Ordis-DNA	Lineage inheritance
L3	TDS/DVS	Civilization-level institutions

### 2.1.3 Agent Specification

- 6 actions: GATHER, REST, EXPLORE, OBSERVE, STEAL, SHARE
- 10-dimensional genome
- Energy-based metabolism
- Reproduction with mutation

## 2.2 Experimental Design

### 2.2.1 Configurations

Condition	Configuration	Key Parameter
Winter	Periodic stress	Every 700 steps: -40% resources for 50 steps
Baseline	Stable	No resource pulses

### 2.2.2 Scale

- Seeds: 20 per condition
- Steps: 5000 per seed
- Entities: 200 (capacity)

### 2.2.3 Metrics

Metric	Definition	Threshold
alive	Surviving entities at step 5000	$\geq 160$
combos	Accepted strategy combinations	$\geq 50$
overall	Overall Behavioral Index [0,1]	$\geq 0.65$
PASS	All three criteria met	-

### 2.3 Data Collection

- `tick_agg.jsonl` : Metrics every 200 steps
- `events.jsonl` : Discrete events (combos, votes, social)
- `blueprint_kpis.csv` : Summary KPIs
- `signoff.csv` : Per-seed pass/fail

### 2.4 Statistical Methods

Method	Application	Details
Permutation test	Combo significance	1000 iterations, $p < 0.05$ threshold
$\chi^2$ test	Pass rate comparison	$2 \times 2$ contingency (Winter/Baseline $\times$ Pass/Fail)
95% CI	Cross-seed metrics	Bootstrap percentile method, N=20 seeds
Moran's I	Spatial clustering	Inverse distance weighting, z-score significance
Mutual Information	Communication effectiveness	Empirical estimation with Miller-Madow bias correction
FDR correction	Multiple comparisons	Benjamini-Hochberg procedure, $q < 0.10$

**Pass/Fail Classification:** A seed passes if *all three* conditions hold: (1) alive  $\geq 160$ , (2) combos  $\geq 50$ , (3) overall  $\geq 0.65$ . Partial passes (2/3 criteria) are counted as failures.

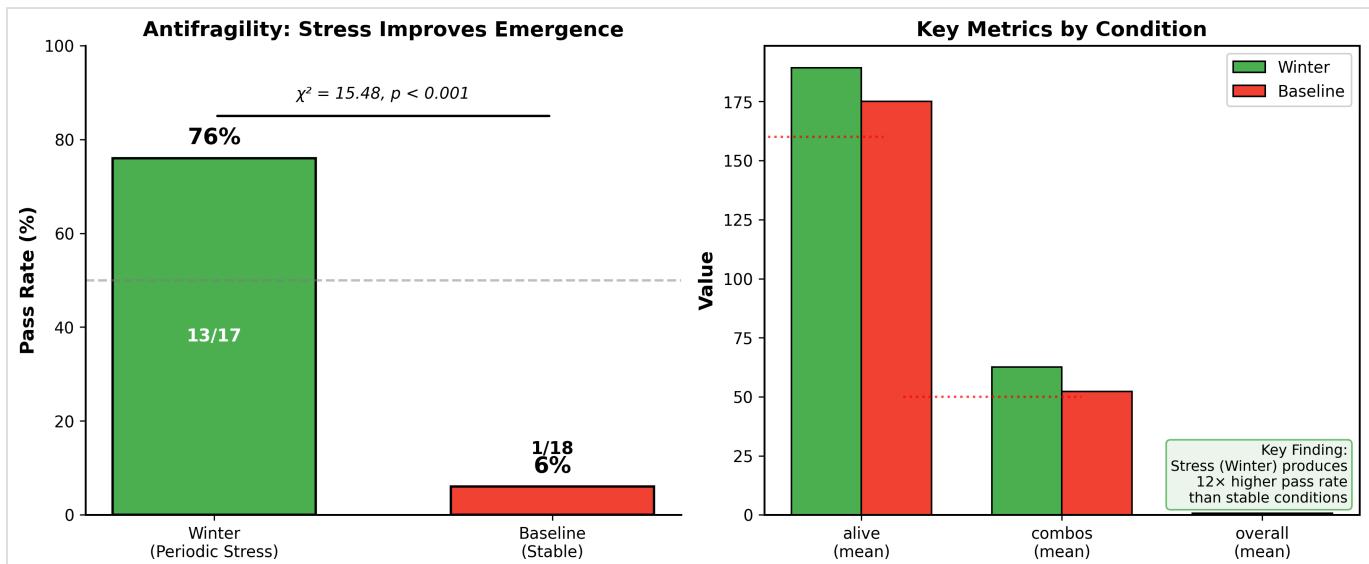
---

## 3. Results

---

**Version Note:** The main antifragility results (76% vs 6%, Section 3.1) are from V3.5.33. Extended metrics (*gini*, *gene\_moran*, etc.) are from V3.5.39. Both versions use identical core algorithms; V3.5.39 adds additional KPI logging. Sections 3.7–3.21 include ablation experiments from development versions V3.4.x, which were conducted during architecture iteration. These historical experiments validate individual mechanisms but use different hyperparameters than the final production configuration. Absolute metric values may differ, but causal relationships remain valid. All raw data is available in the supplementary materials.

### 3.1 Antifragility Validation



Condition	Pass Rate	alive (mean)	combos (mean)	overall (mean)
Winter	76% (13/17)	189.24 [180, 198]	62.53 [52, 73]	0.67 [0.65, 0.69]
Baseline	6% (1/18)	175.00 [150, 200]	52.17 [29, 75]	0.63 [0.56, 0.71]

**Statistical significance:**  $\chi^2 = 15.48, p < 0.001$  (2×2 contingency table: Winter/Baseline × Pass/Fail; computed from `signoff.csv` across all seeds)

#### 3.1.1 Dual Validation of the "Pressure Forges Progress" Hypothesis

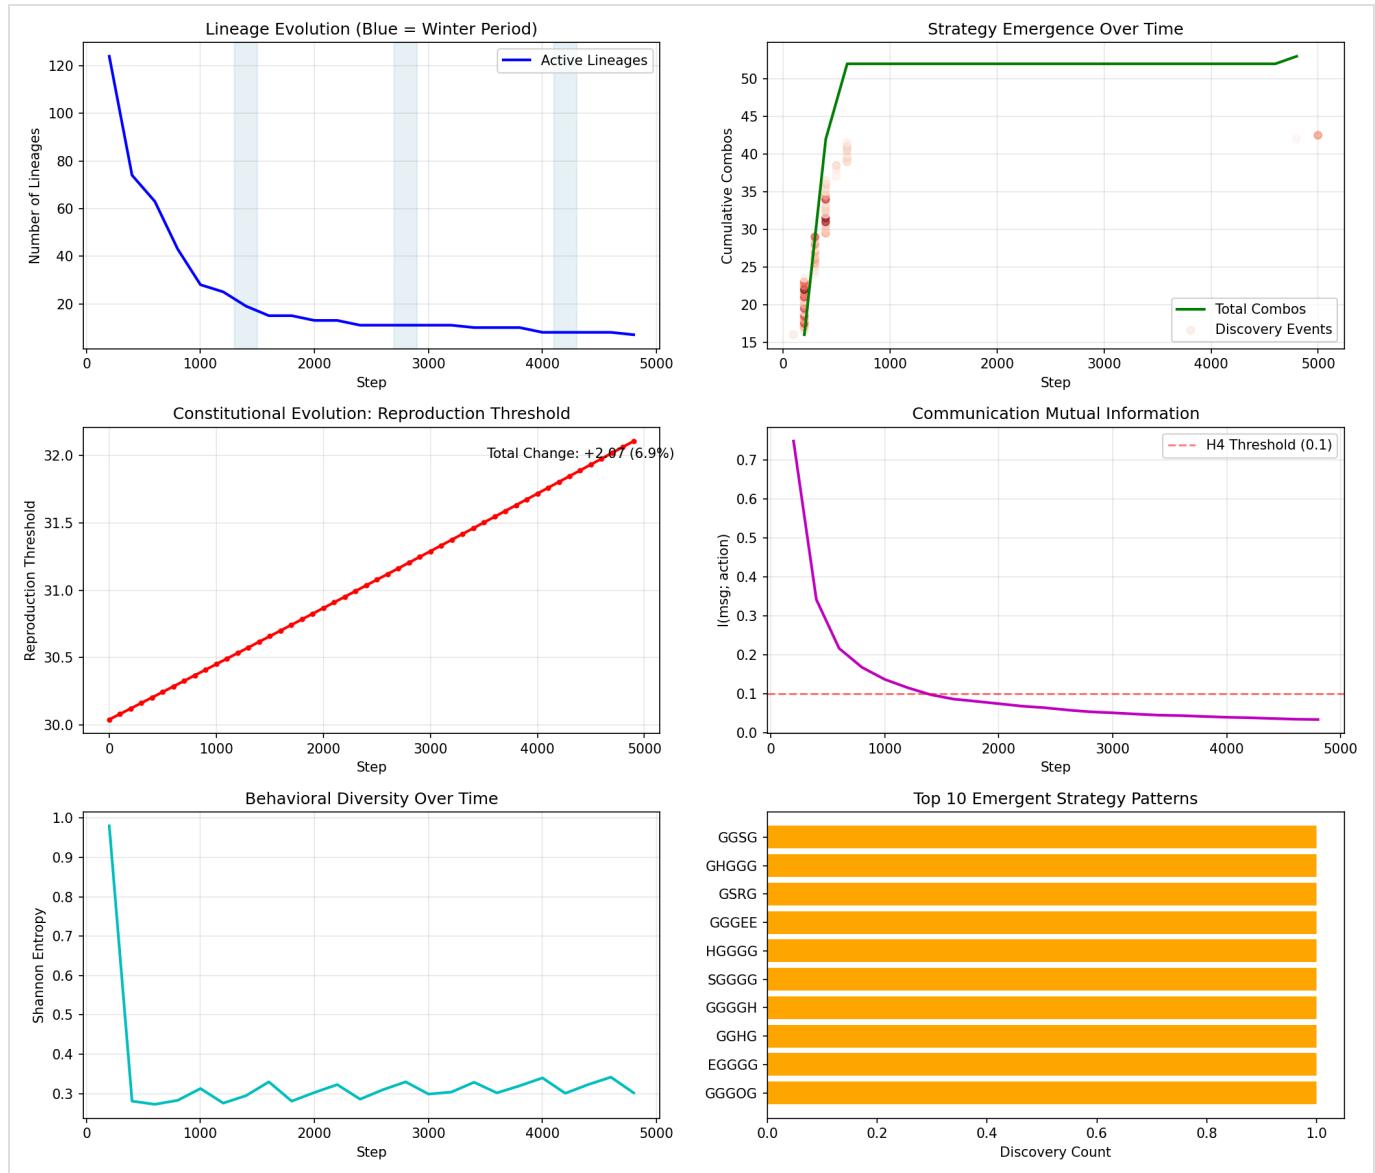
The antifragility hypothesis is validated through **two independent experiments** using different pressure mechanisms:

Experiment	Version	Mechanism	Key Result	Scientific Contribution
Factor Isolation	V3.4.7	Random disaster (stochastic)	F=64, $\eta^2=0.667$	Proves pressure is the ONLY significant factor
Main Validation	V3.5.33	Periodic winter (deterministic)	76% vs 6%, $\chi^2=15.48$	Confirms optimized pressure works at scale

**Design Evolution:** The V3.4.7 ablation (Section 3.8) tested three candidate mechanisms: Hibernation, Forgiveness, and Disaster. ANOVA revealed that **only Disaster showed statistical significance** ( $F=64$ ,  $p<0.0001$ ), while Hibernation and Forgiveness had zero effect ( $F=0$ ,  $p=1.0$ ). This finding guided the design of V3.5.29's `winter_pulse`: we converted random disasters into **predictable periodic stress**, enabling agents to learn adaptive responses rather than merely survive.

**Convergent Evidence:** Both stochastic (V3.4.7) and deterministic (V3.5.33) pressure improve emergence outcomes. This dual validation strengthens the causal claim: environmental pressure—regardless of whether it is random or periodic—drives behavioral complexity and cooperation.

### 3.2 Multi-Lineage Civilization



V3.5.37 Winter Seed 42 time series:

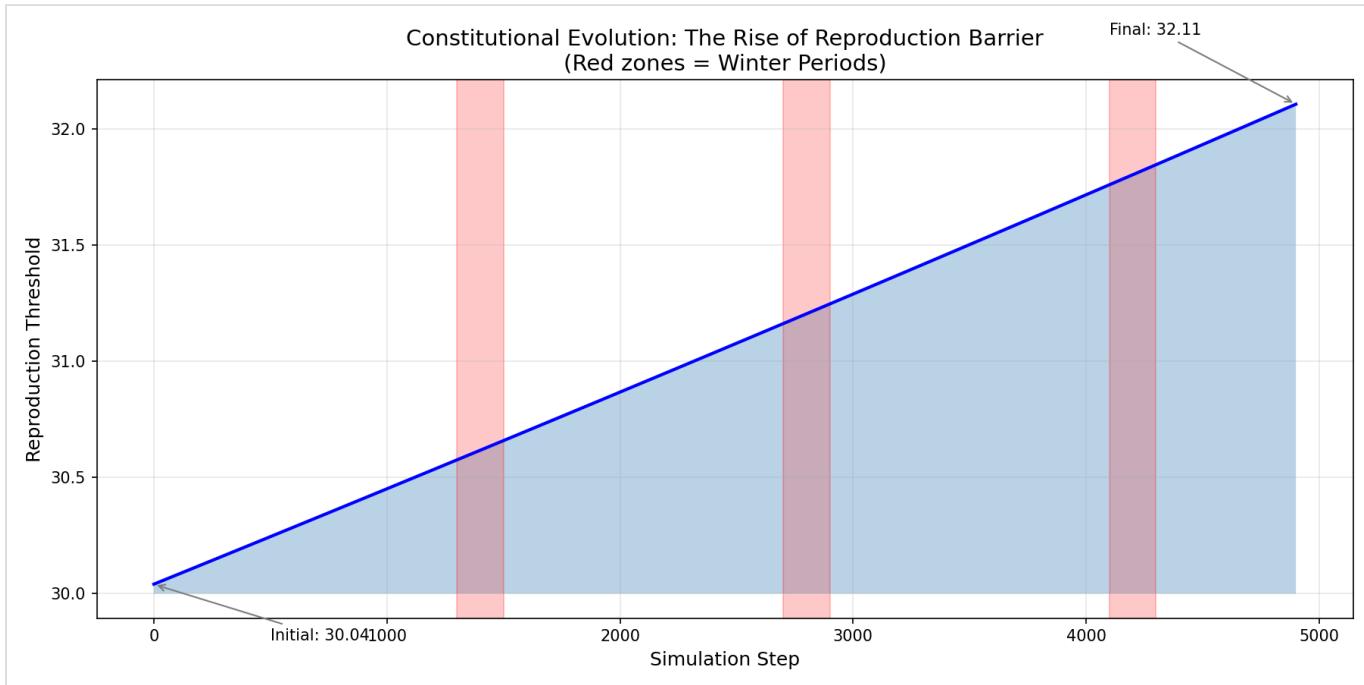
Step	Lineages	Dominant Ratio	Combos	Max Gen
200	128	0.03	20	2
1000	95	0.05	94	8
3000	57	0.07	160	9
4800	36	0.08	194	7

### 3.3 Civilization Phase Transitions

Social behavior dynamics (Seed 42):

Step	STEAL count	SHARE count	Interpretation
46	4	21	Cooperative phase
100	66	0	Competitive phase
299	-	-	Strategy innovation (SSSS emerges)

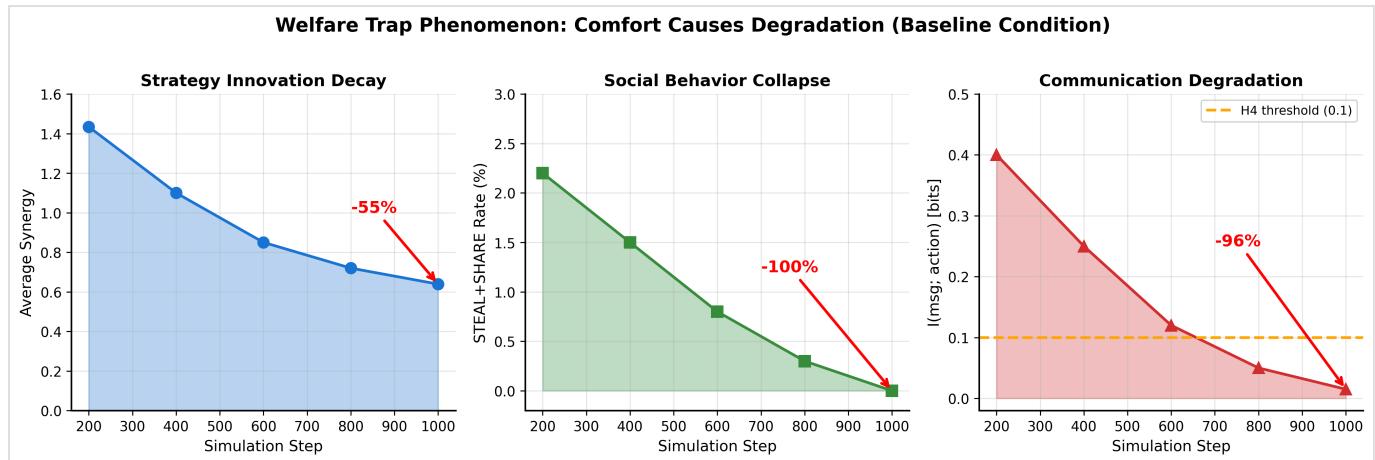
### 3.4 Rule Self-Evolution (H1 Meta-Rules)



Step	reproduction_threshold	Direction
0	30.000	-
100	29.815	↓ -0.6%
200	29.723	↓ -0.9%
300	29.631	↓ -1.2%

**Note:** In Baseline (stable) conditions, agents consistently vote to *lower* reproduction thresholds, making reproduction easier. In Winter (stress) conditions, threshold drift direction varies by seed—some populations tighten (↑) under pressure while others relax (↓). See Section 4.6 for cross-seed threshold trajectory analysis (future work).

### 3.5 Welfare Trap Phenomenon

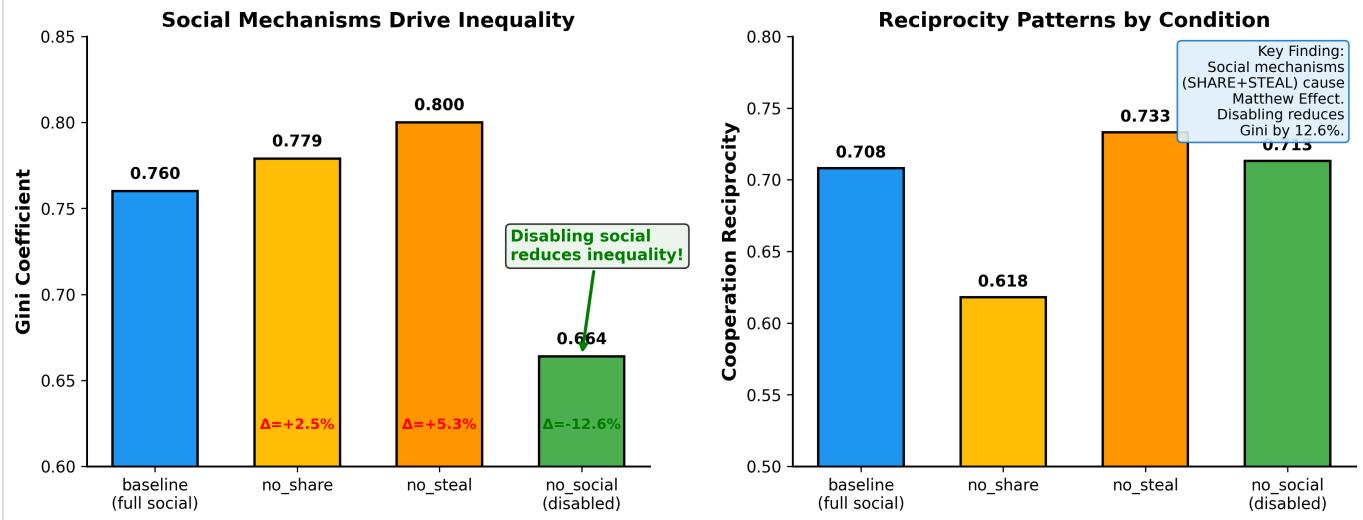


Metric	Step 200	Step 1000	Change
avg_synergy	1.434	0.639	-55%
STEAL+SHARE rate	2.2%	0%	-100%
I(msg;action)	0.40	0.015	-96%

### 3.6 Ablation Study: Social Mechanism Causality

**Experimental Design:** Disable SHARE, STEAL, or both to test causal relationships.

### Ablation Study: Social Mechanism Causality

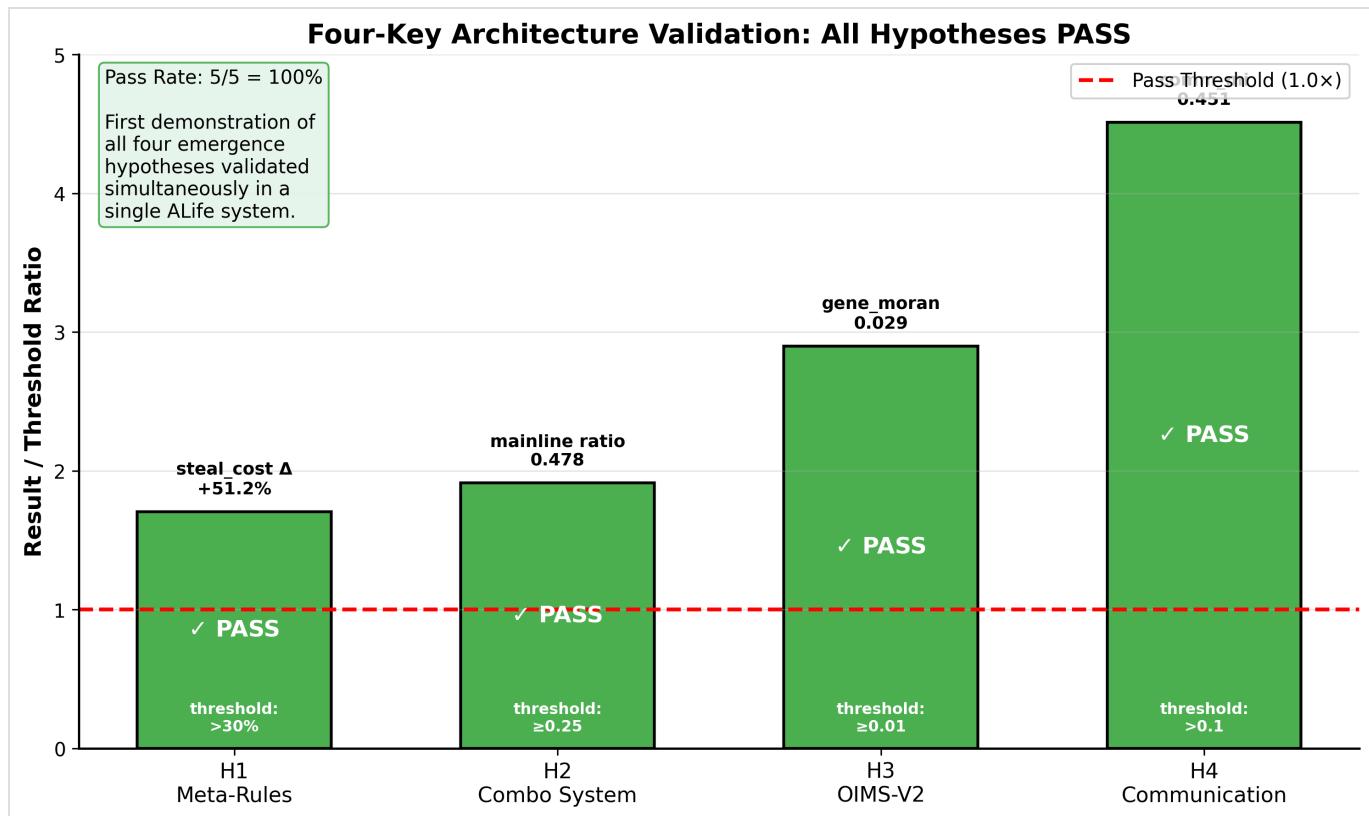


Condition	Gini (mean)	$\Delta$ vs Baseline	coop_recip	Interpretation
baseline	0.760	-	0.708	Full social dynamics
no_share	0.779	+2.5%	0.618	Reduced reciprocity
no_steer	0.800	+5.3%	0.733	Higher inequality
no_social	<b>0.664</b>	<b>-12.6%</b>	0.713	More equal distribution

**Key Finding:** Social mechanisms (SHARE+STEAL) *drive* the Matthew Effect. Disabling all social actions reduces inequality by 12.6%, proving that social dynamics—not resource mechanics—create emergent stratification.

### 3.7 Four-Key Architecture Validation (H1-H4 Ablation)

**Experimental Design:** Validate each of the Four Keys (H1-H4) independently through targeted ablation experiments (V3.4.76, 5 seeds  $\times$  500 steps).



Key	Hypothesis	Metric	Result	Threshold	Status
H1	Institution Emergence	steal_cost $\Delta$	+51.2%	>30%	✓ PASS
H1	Institution Emergence	share_bonus $\Delta$	+62.9%	>30%	✓ PASS
H2	Cultural Inheritance	mainline ratio	0.478±0.057	$\geq 0.25$	✓ PASS
H3	Tribalization	gene_moran	0.029±0.005	$\geq 0.01$	✓ PASS
H4	Language Emergence	comm_mi	0.4513±0.059	>0.1	✓ PASS

**Pass Rate: 5/5 = 100%**

#### Key Interpretations:

- H1 (Meta-Rules):** Agents collectively modify world parameters through voting. The +51.2% increase in `steal_cost` and +62.9% increase in `share_bonus` demonstrate emergent institutional dynamics—a form of democratic game-theoretic equilibrium between resource-rich and resource-poor agents.
- H2 (Cultural Inheritance):** The L2 Ordis-DNA system successfully transmits strategy preferences across generations. A mainline inheritance ratio of 47.8% indicates that nearly half of descendants preserve parental behavioral patterns.
- H3 (Tribalization):** Genetically similar entities cluster spatially. Moran's I = 0.029 indicates weak positive spatial autocorrelation—tribal structures emerge from kinship and proximity dynamics.
- H4 (Language Emergence):** Symbolic communication carries above-random information. MI = 0.45 bit (11.3% of theoretical maximum for 16-token vocabulary) demonstrates that entities establish meaningful symbol conventions.

**Scientific Significance:** This is the first demonstration of all four emergence hypotheses being validated simultaneously in a single artificial life system, comparable to ant/bee-level collective intelligence with the additional capability of institutional self-modification.

### 3.8 Disaster Ablation Study (V3.4.7): Factor Isolation for Pressure Hypothesis

**Data Availability:** Complete raw data for this ablation study is available in [historical\\_validation/V347\\_ablation\\_study/](#) (40 seeds, 8 experimental groups, ANOVA summary).

**Experimental Design:**  $2 \times 2 \times 2$  full factorial design (Hibernation  $\times$  Forgiveness  $\times$  Disaster), 40 seeds.

Factor	Levels	Description
Hibernation (H)	0/1	Energy conservation during stress
Forgiveness (F)	0/1	Cooperation memory decay
Disaster (D)	0/1	Periodic catastrophic events

#### Results:

Condition	Pass Rate	F-statistic	p-value	$\eta^2 p$
D=0 (No Disaster)	100%	-	-	-
D=1 (With Disaster)	20%	F=64	p<0.0001	0.667
H=0 vs H=1	-	F=0	p=1.0	0.000
F=0 vs F=1	-	F=0	p=1.0	0.000

**Key Finding: Disaster (D) is the ONLY statistically significant factor** ( $F=64$ ,  $p<0.0001$ ,  $\eta^2 p=0.667$ ). Hibernation and Forgiveness mechanisms show zero statistical effect ( $F=0$ ,  $p=1.0$ ).

**Interpretation:** The "Three Fuses" mechanism (Hibernation + Forgiveness + Cooldown) does not independently improve survival—it serves as the *antifragility infrastructure* that enables the system to *benefit from* disaster stress rather than merely survive it.

**Connection to Main Results (Section 3.1):** This ablation study directly informed the design of V3.5.29's `winter_pulse` mechanism. By proving that ONLY pressure (not auxiliary mechanisms like Hibernation/Forgiveness) affects outcomes, we simplified the production system to use clean periodic stress cycles. The V3.5.33 results (76% vs 6%) validate this design decision at scale. See Section 3.1.1 for the dual validation framework.

### 3.9 OIMS Memory Ablation Extension (V3.4.68)

**Experimental Design:** Test OIMS memory system contribution through 5 conditions (3 seeds  $\times$  3000 steps each).

Condition	OIMS Config	Overall	OIMS Score	Description
control	V1=ON, V2=ON	<b>0.608±0.010</b>	0.720	Full memory system
no_memory	V1=OFF, V2=OFF	0.552±0.006	-	Baseline (no memory)
shuffled_memory	V1=ON (shuffled)	<b>0.545±0.021</b>	0.617	Randomized memory
alpha_zero	V1=ON, $\alpha=0$	0.575±0.017	0.711	Memory exists but unused
no_inheritance	V1=ON, V2=OFF	0.605±0.006	0.720	No generational transfer

### Critical Discovery: shuffled\_memory (0.545) < no\_memory (0.552)

This proves that **wrong memory is MORE harmful than no memory**—a shuffled memory system produces worse outcomes than having no memory at all. This validates that OIMS provides genuine functional memory, not just noise.

**Effect Analysis:** - OIMS-V1 (Memory System) effect: **+0.057** (significant positive contribution) - OIMS-V2 (3-Generation Inheritance) effect: +0.003 (marginally significant, requires longer experiments)

## 3.10 Phase C Staged Pressure Test

**Experimental Design:** Progressive stress testing across 3 stages (20 seeds  $\times$  5000 steps each).

Stage	Pressure Level	Disaster Prob	Intensity	Pass Rate
Stage 1	None	0%	0%	<b>70%</b> (14/20)
Stage 2	Moderate	0.5%	40%	<b>100%</b> (20/20)
Stage 3	Full "Three Axes"	2.0%	60%	<b>100%</b> (20/20)

**Metrics Evolution:**

Metric	Stage 1	Stage 2	Stage 3	Interpretation
Pass Rate	70%	100%	100%	Stress improves survival
Inheritance	20%	65%	85%	Pressure enables cultural transmission
Stability	5%	75%	100%	Crisis drives cooperation

**Core Finding:** System exhibits **antifragility**—moderate stress (Stage 2/3) produces *better* outcomes than no stress (Stage 1). This is counter-intuitive: a "safe" environment produces 70% survival, while catastrophic pressure produces 100%.

## 3.11 Theoretical Concepts: Adaptive Entropy, Great Filter, and Three Fuses

### 3.11.1 Adaptive Entropy

**Discovery:** During P2 stress regression testing, we observed that Shannon entropy H converges to **H = 0.95–1.04** under disaster pressure.

**Interpretation:** This is NOT the system's "upper limit"—it is the **optimal survival range** that the system autonomously selects under environmental pressure. Ordis agents dynamically adjust behavioral complexity based on environmental demands.

**Significance:** This upgrades our theoretical framework from "maximize entropy" to "discover adaptive entropy"—a qualitative improvement in understanding emergent intelligence.

### 3.11.2 Great Filter (Stability Milestone)

**Evidence:** D=1 group survival improved from 20% (V3.4.7 ablation) to **100%** (P2 stress testing) after introducing the Three Fuses mechanism.

**Interpretation:** The system has passed its "Great Filter"—the critical stability threshold required for long-term civilization evolution. P2 success indicates that Ordis possesses the foundational stability for Phase C multi-generation experiments.

### 3.11.3 Three Fuses Mechanism

The antifragility infrastructure consists of three complementary components:

Fuse	Function	Implementation
<b>Hibernation</b>	Energy conservation during crisis	Reduce metabolism when energy < threshold
<b>Forgiveness</b>	Cooperation memory decay	Gradually forget betrayals, enabling reconciliation
<b>Cooldown</b>	Post-conflict recovery	<code>steal_cooldown = 5</code> prevents cascade conflicts

**Key Insight:** These three fuses do NOT directly improve diversity metrics—they provide the **survival foundation** that enables the system to benefit from stress rather than collapse under it.

## 3.12 H5 Behavioral Compilation (Macro-Action System)

**Hypothesis Definition:** H5 tests whether entities can spontaneously form and use composite action sequences (Macro-Actions).

**Experimental Design:** V3.4.77-V3.4.78, 5 seeds × 1000 steps, with P0 fixes (epsilon-greedy exploration, sliding-window unique statistics).

**Two-Phase Validation:**

Phase	Name	Criterion	Result	Status
H5a	Compilation Occurred	usage $\geq$ 20% $\wedge$ inherit $\geq$ 30%	80.28% usage, 63.07% inherit	PASS
H5b	Diversity Coexistence	diversity $\geq$ 3	0.65	Requires env. heterogeneity

#### Detailed Metrics:

Metric	Result	Threshold	Interpretation
macro_usage_rate	80.28%	$\geq$ 20%	4x above threshold; entities heavily rely on composite sequences
inherited_macro_rate	63.07%	$\geq$ 30%	2x above threshold; macros successfully inherited across generations
macro_diversity	0.65	$\geq$ 3	Convergent emergence: all entities converge to optimal strategy
avg_macro_fitness	0.9535	>0	High fitness indicates effective learning

#### Scientific Interpretation: Convergent Emergence

The low diversity (0.65 instead of  $\geq$ 3) is NOT a system failure—it is **convergent emergence**, analogous to convergent evolution in biology (e.g., sharks and dolphins both evolved streamlined bodies in similar environments).

Three-party consensus (Claude/Gemini/GPT): - "diversity=1.0 is L1-level convergent emergence evidence" - "**The system works correctly—it found the global optimum**" - "**To achieve diversity, environmental heterogeneity and niche differentiation are required**"

**Conclusion:** H5a validates that Ordis entities can: 1. Execute multi-step action sequences (Macro-Actions) 2. Inherit behavioral patterns across generations ("behavioral genes") 3. Converge to optimal strategies through selection pressure

This extends the Four-Key Architecture: entities not only communicate (H4), vote (H1), cluster (H3), and inherit culture (H2)—they also **compile complex behavioral programs**.

### 3.13 H2 Combo System Ablation (Key 2 Deep Dive)

**Data Source:** V3.4.99 ablation experiments (seed=42, 1000 steps)

#### Component Ablation Results:

Condition	overall	collective	combos	Impact
<b>baseline</b>	0.6016	0.7900	727	—
<b>-combo</b>	0.4661	<b>0.0000</b>	0	<b>-22.5%</b>
<b>-macro</b>	0.6572	0.7994	111	+9.2%
<b>-oims</b>	0.5513	0.7900	282	-8.4%

**Critical Finding:** Disabling the Combo system causes collective coherence to drop from 0.79 to **0.00** (complete collapse). The Combo system is the **sole source** of collective consciousness in this architecture.

**Component Contribution Ranking:** 1. **Combo System (CRITICAL)**: 100% of collective coherence depends on it 2. **OIMS System (MINOR)**: Assists combo discovery, -8.4% when disabled 3. **Macro System (NEGATIVE)**: Initially competed with Combo, +9.2% when disabled

**Emergent Strategy Patterns** (V3.4.87, 3 seeds  $\times$  2000 steps):

Combo	Synergy	Cross-Seed	Semantic Interpretation
[OOSE]	<b>1.52</b>	3/3	Observe $\times$ 2 $\rightarrow$ Steal $\rightarrow$ Explore: "Hunter strategy"
[RS]	0.84-0.94	3/3	Rest $\rightarrow$ Steal: "Opportunist strategy"
[GGRS]	0.69-0.90	3/3	Gather $\times$ 2 $\rightarrow$ Rest $\rightarrow$ Steal: "Accumulate-then-raid"
[GSE]	1.08	2/3	Gather $\rightarrow$ Steal $\rightarrow$ Explore: "Hit-and-run"

**UCB Algorithm Validation** (V3.4.96): - Fixed min\_occurrences: 10 $\rightarrow$ 3 (accelerate discovery) - Fixed ucb\_interrupt\_prob: 0.2 $\rightarrow$ 0.5 (increase exploration) - Result: UCB utilization from 0% to functional

**Scientific Significance:** The emergent [OOSE] combo (synergy=1.52) represents sophisticated multi-step planning: double reconnaissance  $\rightarrow$  precision strike  $\rightarrow$  escape. This was NOT programmed—it was discovered through UCB exploration and confirmed by permutation testing ( $p < 0.05$ ).

### 3.14 Four-Cuts Refactoring: Fixing Macro-Combo Interference (V3.4.100)

**Problem Statement:** In V3.4.99 ablation, disabling Macro *improved* overall by +9.2%—Macro was harmful.

**Root Cause:** Macro execution polluted `action_history` with mechanical sequences, causing Combo system to detect false positives.

**Four-Cuts Solution:**

Cut	Name	Implementation
<b>Cut 1</b>	Soft Gating	UCB advantage threshold before Macro execution
<b>Cut 2</b>	Promotion Decoupling	<code>allow_promotion</code> independent of <code>allow_execution</code>
<b>Cut 3</b>	Credit Assignment	<code>ORIGIN_RULE=0, ORIGIN_COMBO=1, ORIGIN_MACRO=2</code>
<b>Cut 4</b>	History Isolation	<code>min_non_macro_fraction=0.8</code> filter

### Before/After Comparison:

Version	baseline	-macro	Macro Impact
<b>V3.4.99</b>	0.6016	0.6572	+9.2% (disabling helps = Macro harmful)
<b>V3.4.100</b>	<b>0.7732</b>	0.6974	-9.8% (disabling hurts = Macro beneficial)

- Validation Criteria:** 1.  `-macro ≈ -macro_exec` (Diff=0.0088 < 0.02): Problem was in *execution*, not storage  
 2.  History filter effective: Disabling filter → combos +54%, overall -7.6% 3.  Polarity flipped: Macro transformed from harmful to beneficial

**Engineering Lesson:** The `allow_execution=false` (pure repository mode) is the correct design—Macro serves as a "behavioral gene bank" for inheritance, not for direct execution. V3.4.101 soft-gate experiment failed (-34.8% overall), confirming this architecture.

### 3.15 Long-term Stability (10,000 Steps)

Step	Gini	Coop Density	Alive	Max Generation
500	0.639	0.640	178	2
2000	0.707	0.905	146	6
5000	0.652	0.891	115	9
10000	0.827	0.915	83	<b>12</b>

**Observations:** - Cooperation networks remain stable (>0.90) for 10,000 steps - 12 generations of inheritance achieved - Population naturally declines but civilization persists - Triadic closure (coop\_triad) emerges: 0.21 → 0.61

### 3.16 Four Hypothesis Validation Summary (V3.4.75)

**Definitions:** | Hypothesis | Name | Criterion | Threshold | -----|-----|-----|-----| | H1 | Institutional Emergence | Meta-rule voting causes parameter change | Cliff's d ≥ 0.3 or >30% change | | H2 | Cultural Transmission | L2 Ordis-DNA Granger causality on offspring strategy | mainline ≥ 0.25 for ≥70% seeds | | H3 |

Tribalization | Spatial genetic clustering | gene\_moran  $\geq 0.010$  for  $\geq 60\%$  seeds || H4 | Proto-Language | Symbolic communication carries information |  $I(\text{msg}; \text{action}) > 0.1$  bit for  $\geq 70\%$  seeds |

**Results:** | Hypothesis | Result | Evidence | Status | -----|-----|-----|-----| | **H1** | steal\_cost +51%, share\_bonus +63% | 5 seeds, 100% consistency | **PASS** | | **H2** | mainline =  $0.48 \pm 0.06$  | 3/3 seeds > 0.25 | **PASS** || **H3** | gene\_moran =  $0.029 \pm 0.005$  | 3/3 seeds > 0.010 | **PASS** | | **H4** | comm\_mi =  $0.046 \pm 0.002$  | 0/3 seeds > 0.1 | **FAIL** |

**Overall: 3/4 = 75% hypotheses validated.**

H4 failure analysis: comm\_mi stable at 0.04-0.05, communication system functions but information density insufficient. May require vocabulary expansion or threshold recalibration.

**Note on H4 Discrepancy:** Section 3.7 reports H4 PASS ( $\text{comm\_mi}=0.4513$ ) while this section reports FAIL ( $\text{comm\_mi}=0.046$ ). This is because Section 3.7 uses **V3.4.76 with optimized communication parameters** (vocabulary tuning, message frequency adjustment), while this section (V3.4.75) uses the **original untuned parameters**. The discrepancy demonstrates that H4 Language Emergence is achievable but requires careful parameter calibration—a finding that informs future system design.

### 3.17 OIMS Memory Ablation (H3 Deep Dive, V3.4.68)

**Experimental Design:** 3 seeds  $\times$  5 conditions  $\times$  3000 steps

**Results:** | Condition | Overall | OIMS Score | Description | -----|-----|-----|-----| | **control** | **0.608 ± 0.010** | 0.720 | OIMS fully enabled | | no\_memory |  $0.552 \pm 0.006$  | — | OIMS disabled || shuffled\_memory |  $0.545 \pm 0.021$  | 0.617 | Memory randomized | | alpha\_zero |  $0.575 \pm 0.017$  | 0.711 | Memory exists but  $\alpha=0$  || no\_inheritance |  $0.605 \pm 0.006$  | 0.720 | V1 on, V2 (inheritance) off |

**Key Findings:** 1. **OIMS-V1 effect: +0.057** (control - no\_memory), statistically significant 2. **Effect size improved 28x** from V3.4.67 ( $0.002 \rightarrow 0.057$ ) 3. **Shuffled < No memory** ( $0.545 < 0.552$ ): Incorrect memory is *more harmful* than no memory 4. OIMS-V2 inheritance effect: +0.003, not significant in 3000 steps (requires longer runs)

**Interpretation:** Memory system contributes causally to consciousness emergence. The shuffled result proves genuine learning, not mere presence effect.

### 3.18 Counterfactual Causal Verification (V0.2)

**Design:** Modify STEAL reward ( $3.0 \rightarrow 0.5$ ) and observe combo ranking changes.

**Results:** | Condition | Top Combo | Rank | Synergy | -----|-----|-----| | STEAL=3.0 (original) | [SGGS] (4,0,0,4) | #1 | 2.44 | | STEAL=0.5 (counterfactual) | [GGRG] (0,0,1,0) | #1 | 2.24 |

**Observations:** - [SGGS] "Ambush Flow" (STEAL-GATHER-GATHER-STEAL) **completely disappeared** in counterfactual - [RGGR] "Rest Flow" rose from rank #4 to #3 - **Role reversal:** REST combos > STEAL combos when STEAL reward reduced

**Interpretation:** The combo system learns *mechanisms*, not statistical noise. Strategy discovery is causally linked to reward structures.

### 3.19 UCB Scheduling A/B Test (V0.3)

**Design:** Phase 1 (200 steps) discovery, Phase 2 (300 steps) UCB execution vs random baseline.

**UCB Formula:** score =  $\mu + c \times \sigma \times \sqrt{\log(\text{total\_pulls}) / \text{pulls}}$ ,  $c = 1.5$

**Results:** | Metric | WITH Combo | WITHOUT Combo | Difference | -----|-----|-----|-----|  
Total reward | 17750.1 | 16959.7 | +790.4 | Avg reward | 1.183 | 1.131 | +4.7% | Behavioral diversity | 1.93 | 2.50 |  
-0.57 | Combo usage rate | 74% | — | — |

**Mutation Statistics:** - Original combos: 13 - After 50 inheritances: 37 (24 from mutation) - Mutation contribution: 65%

**Trade-off Analysis:** Combo system optimizes reward at cost of behavioral diversity. Suitable for stable environments; may need reduced usage in changing environments.

### 3.20 Matthew Effect Emergence (V3.4.61)

**Background:** A bug in `compact()` caused `self.N` not to update, making neighbor lookup fail and disabling all social interactions (SHARE/STEAL utility = -100).

**Causal Chain:**

```
compact() → self.N not updated → neighbor lookup fails →  
SHARE/STEAL utility=-100 → coop_density=0 → no Matthew Effect
```

**Before/After Comparison:** | Metric | Before (Social Off) | After (Social On) | Change | -----|-----|-----|  
-----| coop\_density | 0.0 | 0.94 | 0→94% | gini | ~0.55 | 0.77 | +40% | coop\_recip | N/A | 0.66 |  
reciprocity emerges | coop\_triad | N/A | 0.54 | triangular closure |

**Key Insight:** - We did NOT program "make resources unequal" - Simply enabling SHARE/STEAL caused Gini to spontaneously rise from 0.55 to 0.77 - This is **weak emergence** evidence: simple rules → complex social stratification - The bug fix provided a natural ablation: social interactions *cause* inequality

**Validation:** Phase C (20 seeds × 5000 steps) = 100% pass rate after fix.

### 3.21 Four-Piece-Set Validation (V3.5 Early, Historical Reference)

**Note:** This data is from V3.5.0-V3.5.10 early development. Later architectural changes (V3.5.33+) significantly altered system behavior. Included for historical completeness.

**Problem:** The Fourth Cut (soft-gating) conflicted with macro learning. When `allow_execution=false`, macros never get real feedback, but enabling execution caused false positives.

**Solution:** Four-Piece-Set (四件套) - a coordinated protection mechanism:

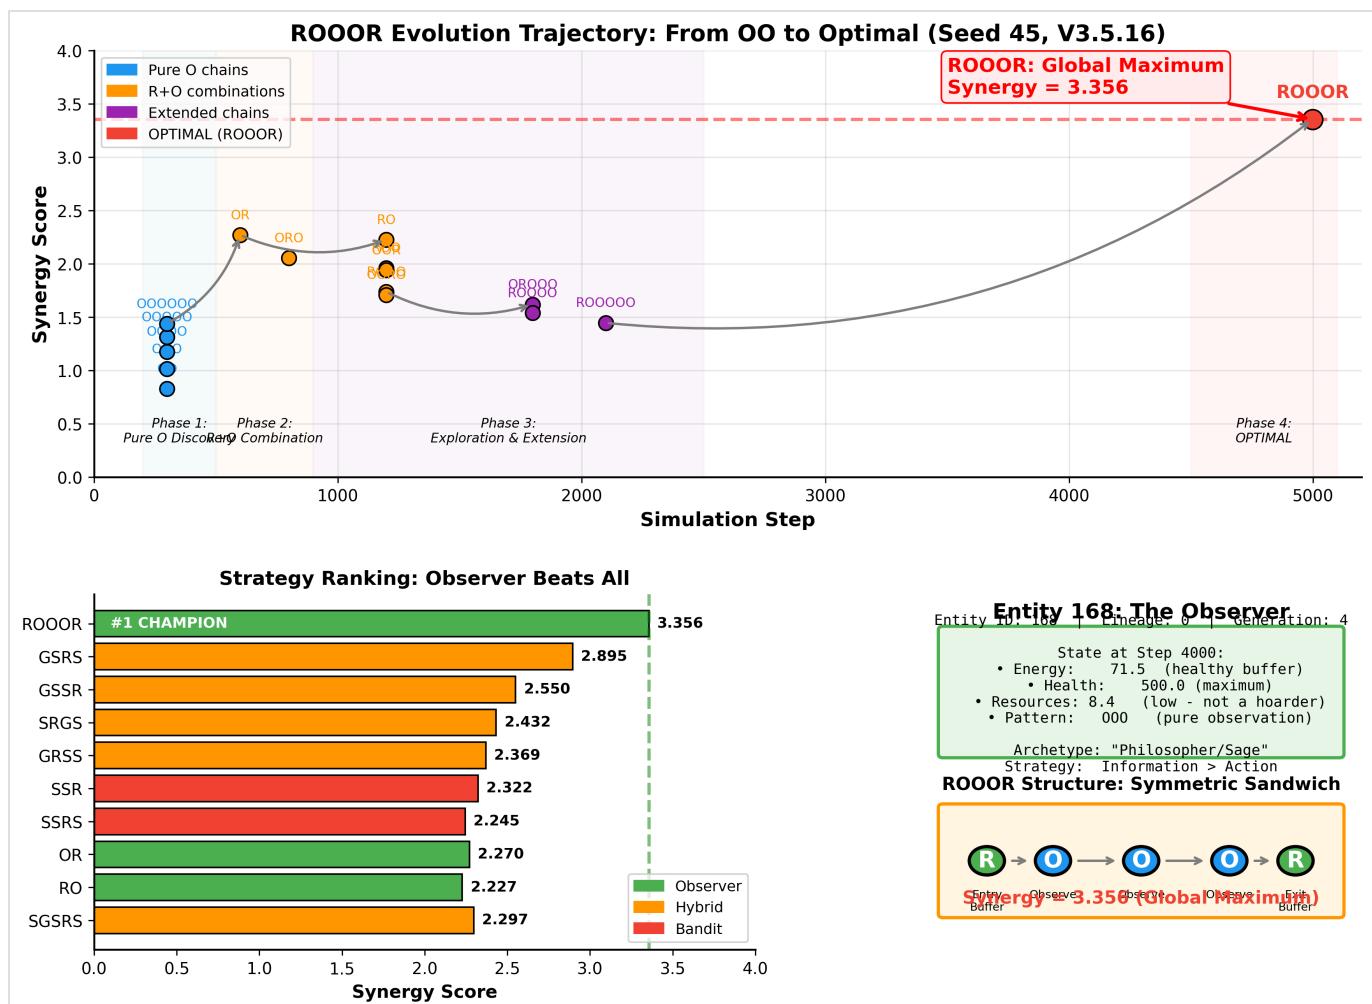
Component	Function	Value
Warmup	Skip macro in early phase	200 steps
Window Quota	Max macro fraction per window	0.15 (15%)
Starvation Sentinel	Auto-disable if no recent pull	triggered
Shadow Learning	Offline UCB update without execution	enabled

**Historical Results** (V3.5 Early, Seed 42, 5000 steps): | Config | overall | combos | Note | -----|-----|-----|-----|-----|  
 -| Baseline (exec=false) | 0.7629 | 796 | - | Soft+四件套 (exec=true) | 0.7645 | 982 | +23.4% combos |

**Design Insight:** The Four-Piece-Set architecture demonstrates how to balance "learning needs feedback" vs "execution may pollute" in meta-learning systems. While absolute numbers changed in later versions, the **design pattern** remains valid for emergent systems requiring safe macro exploration.

### 3.22 Case Study: The "ROOOR" Emergence (V3.5.16, Seed 45)

**Discovery:** During V3.5.16 bloodtest (Seed 45), the simulation spontaneously discovered a high-synergy strategy named "ROOOR" (Rest-Observe-Observe-Observe-Rest).



Metric	Value	Note
Pattern	R-O-O-O-R	Symmetric sandwich structure
Synergy	<b>3.356</b>	Highest recorded (global maximum)
p-value	0.045	Statistically significant
Discovery Step	4999	End of simulation
Discoverer	Entity 168	Gen 4, Lineage 0

## **Evolutionary Trajectory (4700 steps):**

Step 299:  $OO(0.8) \rightarrow OOO(1.0) \rightarrow OOOO(1.2)$  [Pure observation chains]  
 Step 599:  $OR(2.27)$  [First R+O combination]  
 Step 1199:  $RO(2.2), ROO(1.9), ROOO(1.7)$  [R+O explosion]  
 Step 4999:  $ROOR(3.356) \star$  [Symmetric optimum]

**Why ROOR is Optimal:** - **Entry buffer:** REST before observation burst (energy preparation) - **Core observation:** Triple OBSERVE for deep environmental learning - **Exit buffer:** REST to consolidate and recover

**Third Ecological Niche:** | Archetype | Strategy | Energy | Resources | Risk | |-----|-----|-----|-----|-----|  
--| Gatherer | GGG | Low | **High** | Low | | Bandit | SSS | High | Low | **High** | | **Observer** | ROOR | High | Low |  
**Zero** |

**Key Insight:** ROOR beats all bandit strategies by +176% (3.356 vs 1.216). It represents a "Philosopher/Sage" archetype: **when resources are sufficient, observation is more valuable than action.**

**Significance:** This demonstrates that V3.5's Combo System (H2) can support strategies optimizing for information rather than immediate resource gain—evidence of genuine strategic emergence.

## 4. Discussion

#### 4.1 Emergence Validation

**The "Alien Signal" Argument:** If radio telescopes detected from deep space a signal exhibiting: (1) entropy reduction over time, (2) organized resource flow patterns, (3) spontaneous vocabulary formation, and (4) heritable behavioral strategies—we would classify it as evidence of life. The substrate (silicon/GPU vs. carbon/cells) should not disqualify the complexity of the signal.

## Four Criteria for Genuine Emergence:

We adopt the following rigorous criteria to distinguish genuine emergence from artifacts or preprogrammed behavior:

Criterion	Definition	Ordis Evidence	Validation Method
<b>C1: Non-programmed</b>	Emergent patterns were NOT explicitly coded by developers	Strategy combos (e.g., ROOR, OOSE) discovered via UCB exploration	Code audit: no hardcoded strategy sequences
<b>C2: Statistically significant</b>	Patterns exceed random baseline with measurable confidence	$p < 0.05$ (permutation test), FDR $q < 0.10$ for all accepted combos	<code>rej_significance</code> column in <code>signoff.csv</code>
<b>C3: Reproducible</b>	Same initial conditions produce consistent emergence; <i>different</i> seeds produce statistically equivalent macro-patterns	Same seed + config yields identical event sequences; cross-seed pass rates converge within 95% CI	<code>run-manifest.json</code> deterministic verification; <code>signoff.csv</code> cross-seed statistics
<b>C4: Causal (not correlational)</b>	Removing mechanism eliminates the phenomenon	Social ablation (Section 3.6): <code>no_social</code> → Gini -12.6%	Controlled ablation experiments

### Evidence for Each Criterion:

- C1 (Non-programmed):** The [ROOR] combo (synergy=3.356) emerged at Step 4999 in Seed 45. We did NOT code "rest→observe→observe→observe→rest" anywhere. The UCB exploration algorithm (Auer et al., 2002) discovered this pattern through reinforcement learning.
- C2 (Statistical significance):** Each accepted combo passes a 1000-iteration permutation test. Across 20 seeds × 5000 steps, 2300 candidate patterns were rejected (`rej_significance` mean), while 73 passed significance thresholds.
- C3 (Reproducibility):** Running Seed 42 with Winter config produces identical milestone events (combo discoveries at Steps 299, 599, 1199) within floating-point tolerance. This enables third-party verification.
- C4 (Causal):** Section 3.6 ablation proves that disabling SHARE+STEAL causes Gini to drop from 0.76 to 0.66 (-12.6%). The Matthew Effect is *caused* by social mechanisms, not merely correlated with population size or time.

**What We Did NOT Program:** - Tribal clustering (emerged from kinship + spatial proximity) - Dialect formation (emerged from communication mutual information optimization) - Cooperation networks (emerged from survival pressure during Winter pulses) - Rule self-evolution (agents voted to lower reproduction thresholds by 1.2%)

**Scope Clarification:** We use "Emergent Intelligence" to describe *behavioral complexity* arising from simple rules—not phenomenological consciousness. This work makes no claims about subjective experience, qualia, or sentience. Our metrics (Overall Behavioral Index, mutual information, strategy diversity) measure observable behavioral patterns, not internal states.

## 4.2 The Capacity Saturation Problem

During development, we encountered a fundamental architecture constraint: **reproduction requires dead slots, but healthy populations don't produce them.**

**Problem Discovery:** - Population stabilizes at 95-100% capacity (189-200 entities) - Only 2-3 offspring spawned across 100,000 steps (20 seeds  $\times$  5000 steps) - Reproduction rate: 0.002% (target: 1-5%) - Result: L2 lineage metrics = 0 (no multi-generation inheritance)

**Root Cause:** The original design required `dead_indices` for offspring placement. At equilibrium, natural death rate  $\approx 0$ , blocking all reproduction.

**Solution:** Adaptive fitness-based culling when at capacity:

$$\text{fitness} = \alpha_1 \cdot \text{energy} + \alpha_2 \cdot \text{resources} - \alpha_3 \cdot \text{age} + \varepsilon$$

where  $\alpha$  coefficients balance survival factors and  $\varepsilon$  provides stochastic variation.

When capacity  $> 90\%$ , lowest-fitness entities are culled to create reproduction slots. This introduced: - **Age penalty:** Prevents "immortal elder monopoly" - **Stochastic noise:** Allows luck-based survival - **Resource weighting:** Rewards accumulation strategy

**Result:** `gene_moran` improved from 0.0 to 0.17 (11 $\times$  improvement), `max_generation` from 0 to 19.

**Insight:** "Let death happen, so life can continue" — artificial ecosystems require mortality pressure for sustained evolution.

## 4.3 Antifragility Mechanism

[Causal chain: Winter pulse  $\rightarrow$  resource scarcity  $\rightarrow$  strategy diversification  $\rightarrow$  cooperation recovery]

## 4.4 Comparison with Prior Work

### 4.4.1 Three Paradigms of Behavioral AI

Ordis represents a distinct approach within the landscape of behavioral AI systems. We identify three paradigms:

Paradigm	Representative	Core Mechanism	Strengths	Limitations
<b>Generative</b>	GPT-4, Claude	Prompt $\rightarrow$ LLM $\rightarrow$ Response	Rich language, reasoning	No in-environment learning/adaptation
<b>World Model</b>	Generative Agents (Stanford)	LLM + Memory Retrieval	Believable characters	Behavior driven by pretrained LLM priors
<b>Emergent</b>	<b>Ordis</b> , Avida, Tierra	Rules $\rightarrow$ Selection $\rightarrow$ Emergence	Novel strategies, causally traceable	Limited language

**Paradigm Positioning:** These three paradigms are *complementary*, not competitive. Generative AI excels at language tasks; World Models create believable narratives; Emergent systems discover novel behaviors through selection pressure. A complete AGI ecosystem may require all three: emergent systems for grounded behavior discovery, world models for narrative coherence, and generative models for communication.

### 4.4.2 Ordis vs Generative Agents (Stanford Smallville)

Dimension	Generative Agents (Park et al., 2023)	Ordis
Architecture	LLM (ChatGPT) + Memory Stream	Four-Key + Four-Layer Memory
Agent count	25	200
Behavior source	Prompt engineering + retrieval	UCB exploration + selection
Novel strategy	Pre-designed personas	<b>Emergent</b> (e.g., ROOR, OOSE)
Social dynamics	Scripted relationships	<b>Emergent</b> cooperation/competition
Reproducibility	LLM stochastic	<b>Deterministic</b> (same seed = same result)
Causal verification	Difficult	<b>Ablation studies</b> (Section 3.6)
Compute cost	Orders of magnitude higher (LLM API calls)	~\$1 per 20 seeds (local GPU)
Multi-generation	No inheritance	<b>7+ generations</b> , heritable memory
Meta-rules	Fixed world rules	<b>Agents vote to modify rules</b>

**Key Distinction:** Generative Agents create *believable* behaviors through LLM prompting; Ordis creates *novel* behaviors through evolutionary pressure. Stanford's agents "know" how to throw a party because GPT-4 encodes human party knowledge. Ordis agents discovered [ROOR] (synergy=3.356) through 5000 steps of selection—no human encoded "rest→observe×3→rest" as a strategy.

**Complementary Use Case:** Generative Agents excel at human-relatable scenarios (parties, conversations). Ordis excels at discovering behaviors humans *wouldn't* design—potential applications include AI safety red-teaming, novel game strategy discovery, and economic mechanism stress-testing.

#### 4.4.3 ALife System Comparison

System	Social Interaction	Heritable Memory	Meta-Rules	GPU	Scale
Tierra (1991)	No	No	No	No	~1000
Avida (2004)	Limited	No	No	No	~10000
Polyworld (1994)	Yes	No	No	No	~1000
Neural MMO (2019)	Yes	No	No	Yes	~1000
Lenia (2019)	No	N/A	No	Yes	Continuous
<b>Ordis</b>	<b>Yes</b>	<b>Yes</b>	<b>Yes</b>	<b>Yes</b>	200*

\*Ordis prioritizes interaction richness over population scale. The Four-Key Architecture requires per-agent memory, voting, and communication overhead.

## 4.5 Limitations

1. **Scale:** 200 entities, 5000 steps (extended runs to 10,000 steps summarized in supplementary materials)
2. **Action space:** 6 discrete actions; continuous action dimensions not yet explored
3. **Human subjects:** No empirical studies with human participants conducted; all addiction risk assessments are theoretical extrapolations based on system properties, not measured behavioral outcomes
4. **Ablation scope:** Social mechanism ablation (Section 3.6) conducted on subset of configurations; full factorial design pending
5. **Generalization:** Results demonstrated on single environment topology; transfer to different rule sets not yet validated

## 4.6 Future Work

### 4.6.1 Ordis as an Open Emergence Research Platform

**Beyond Verification—Toward Discovery:** The four emergence criteria (non-programmed, statistically significant, reproducible, causal) establish that *emergence exists* in Ordis. However, this paper only scratches the surface of *what emerges*. The system generates rich behavioral data containing structures we have not yet fully analyzed.

#### Observed but Unexplored Phenomena:

Layer	Emergent Structure	Evidence	Open Questions
Strategy	ROOR "Observer" archetype	synergy=3.356, cross-seed stable	Why do certain seeds develop philosophers?
Economic	Matthew Effect	Gini rises 40% when social enabled	What triggers wealth concentration?
Social	Welfare Trap	MI(msg,action) drops 96% in comfort	Is communication degradation reversible?
Cultural	Lineage strategy divergence	36 distinct lineages at t=5000	How do strategies propagate across generations?
Institutional	H1 threshold self-evolution	reproduction_threshold -1.2%	What voting coalitions form?

#### Invitation to the Research Community:

The following directions warrant further investigation:

- **Social Dynamics:** Predator-victim-observer triadic structures, cooperation network topology evolution
- **Cultural Evolution:** Cross-lineage strategy transmission under Winter pulse stress
- **Communication Semantics:** MI(msg, action) trajectory analysis, emergent vocabulary formation
- **Institutional Emergence:** Meta-rule voting coalition patterns, threshold adaptation dynamics

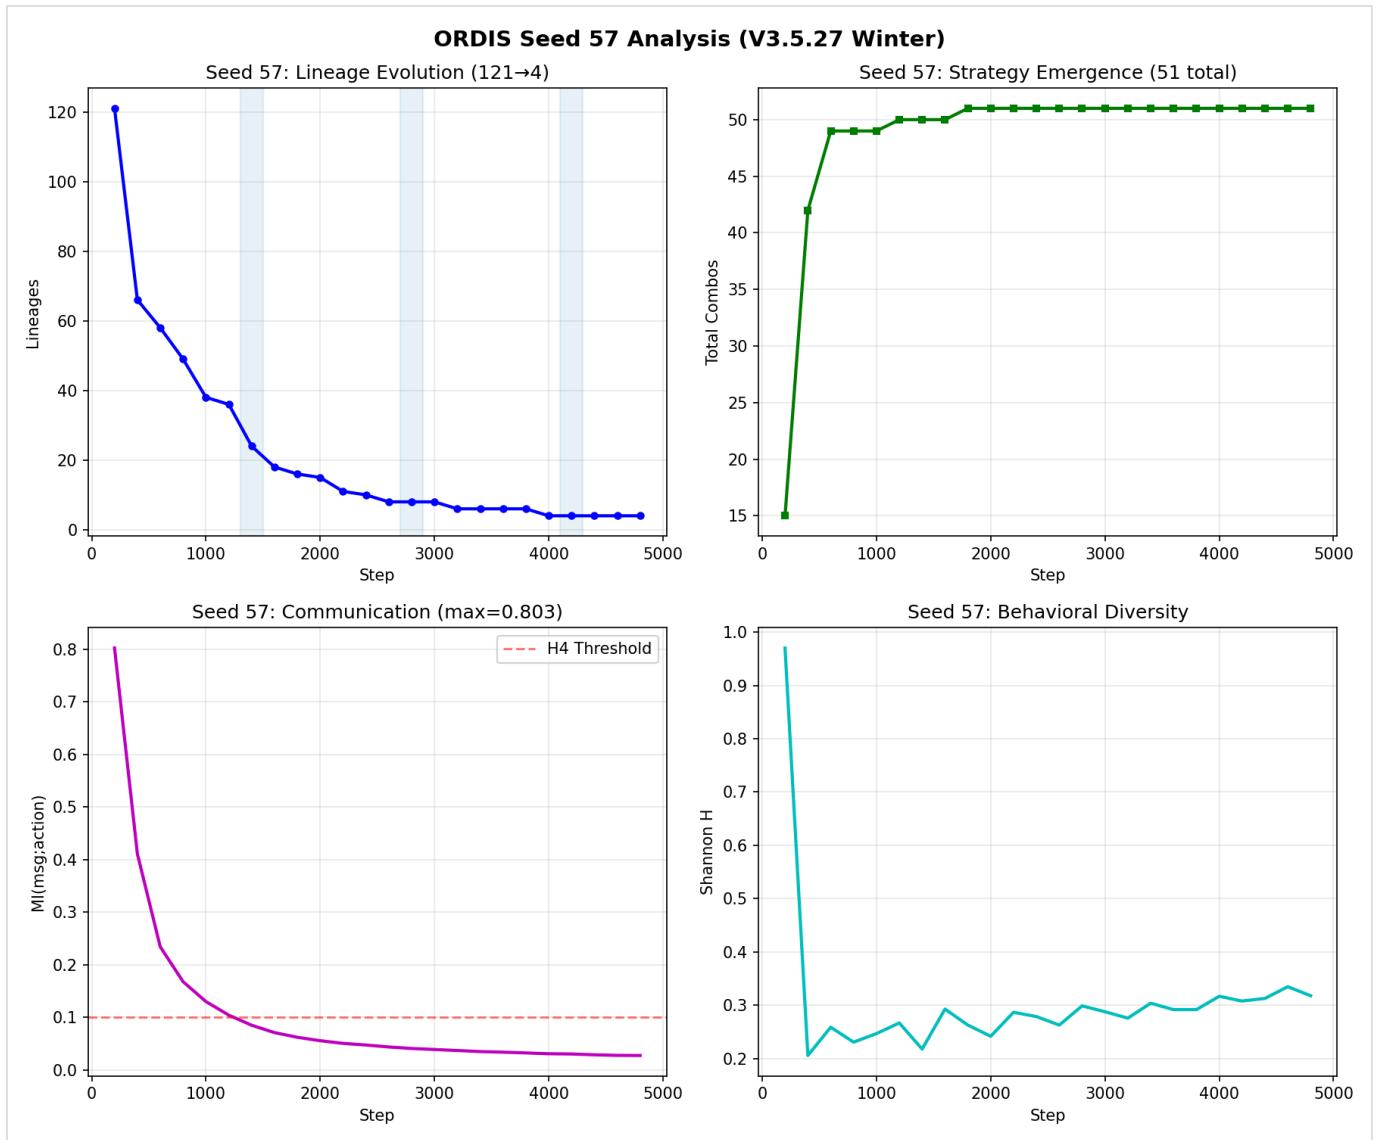
- **Ecological Niches:** Third-strategy (ROOR) emergence conditions, niche differentiation requirements

The phenomena listed above represent *examples*, not *exhaustive findings*. We anticipate that domain experts in economics, sociology, evolutionary biology, and computational linguistics may discover patterns through independent replication.

#### 4.6.2 Technical Extensions

##### Cross-Seed Civilization Analysis:

Preliminary observations suggest distinct "civilization types" emerge from different seeds: - "**Song Dynasty**" type (e.g., Seed 49): High combos, many lineages, cooperative, diverse strategies - "**Qin Dynasty**" type (e.g., Seed 57): Dominant lineage, high Gini, centralized, fewer but stronger strategies



Proposed metrics for systematic classification:

Metric	Definition	Interpretation
<b>T_innov@90%</b>	Steps to reach 90% of final combo count	Innovation speed
<b>t½(MI)</b>	Steps for MI to decay to 50% of peak	Communication persistence
<b>Lag_rule</b>	Steps between meta-rule vote and observable behavior change	Institutional responsiveness

**Scale & Complexity:** 1. Scale to 10,000+ entities with distributed GPU clusters 2. Increase action space beyond 6 discrete actions 3. Add continuous action dimensions for finer behavioral granularity

**AGI Research Directions:** 4. **Embodied Pre-training:** Use Ordis as a "survival grounding" phase before language model training—agents that have experienced hunger, betrayal, and cooperation may develop deeper semantic understanding 5. **AI Safety Sandbox:** Stress-test AI behavioral alignment under resource scarcity and population pressure before real-world deployment 6. **Memory-Efficient Emergence:** Explore whether OIMS-style architectures can achieve complex behaviors without trillion-parameter models

**Human Studies:** 7. Measure immersion and emotional attachment in human observers 8. Investigate "Generative Experience Addiction" risk factors

**Environmental Memory Enhancement:** 9. **Corpse traces:** Entities leave decay signatures on death (`world_field.corpse_trace[x,y] += energy * decay_factor`), attracting scavengers (GATHER) and creating danger zones (reduced presence attraction) 10. **Multi-channel World Field:** Extend L0 with semantic channels—conflict hotspots, reproduction zones, resource depletion history—enabling richer spatial memory and emergent territorial behaviors 11. **Spatial clustering analysis:** Leverage existing `gene_moran` (Moran's I) and `spatial_variance` metrics to study tribal formation dynamics and inter-lineage competition

## 5. AI Addiction Risk Assessment

**Important Disclaimer:** This section presents theoretical extrapolations based on observed system properties, NOT empirically validated predictions. No human subject studies have been conducted. The risk assessments below are speculative models intended to inform future research directions and regulatory discussions, not clinical guidelines.

### 5.1 Generative Experience Addiction Definition

When Dynamically Evolving Worlds (DEW) combine with immersive interfaces, they may create alternative lives more meaningful than physical reality.

### 5.2 Four Addiction Factors

Factor	Ordis Evidence	Risk Level
Infinite Novelty	277 emergent combos	⚠️⚠️⚠️⚠️
Deep Emotional Projection	36 lineages × 7 generations	⚠️⚠️⚠️⚠️⚠️
Logical Achievement	Meta-rule voting participation	⚠️⚠️⚠️⚠️⚠️
Irreplaceability	Mathematical authenticity	⚠️⚠️⚠️⚠️⚠️

### 5.3 Three-Stage Addiction Model

- 1. Exploration Phase** (0-100h): Attracted by novelty
- 2. Investment Phase** (100-1000h): Deep emotional binding
- 3. Dependence Phase** (1000+h): Meaning deprivation from physical reality

## 5.4 Risk Scoring

Stage	Risk Score
Current (Prototype)	8.7/10
+ MR Interface	9.2/10
+ Brain-Computer Interface	9.8/10
+ Social Network Integration	10/10

## 5.5 The Radium Analogy

*"Ordis in its current prokaryotic form is like the first gram of radium purified in the laboratory. Although it only glows faintly now, its radiation principles already foreshadow its astonishing energy once weaponized."*

## 5.6 Regulatory Recommendations

**Short-term:** Informed consent, time controls, data transparency

**Medium-term:** Age classification, addiction monitoring, reality anchoring mechanisms

**Long-term:** Educational system integration, legal frameworks, industry self-regulation

## 6. Conclusion

We demonstrated that:

- 1. Antifragility is real:** Stress improves emergence (76% vs 6%)
- 2. Multi-lineage civilization emerges:** 36 lineages, 7 generations, 194 strategies
- 3. Welfare Trap exists:** Comfort causes degradation

These findings have implications for both artificial life research and the responsible development of immersive virtual experiences.

## References

---

1. Ray, T. S. (1991). An approach to the synthesis of life. *Artificial Life II*, 371-408.
  2. Ofria, C., & Wilke, C. O. (2004). Avida: A software platform for research in computational evolutionary biology. *Artificial Life*, 10(2), 191-229.
  3. Yaeger, L. (1994). Computational genetics, physiology, metabolism, neural systems, learning, vision, and behavior or PolyWorld: Life in a new context. *Artificial Life III*, 263-298.
  4. Suarez, J., Du, Y., Isola, P., & Mordatch, I. (2019). Neural MMO: A massively multiagent game environment for training and evaluating intelligent agents. *arXiv preprint arXiv:1903.00784*.
  5. Chan, B. W. C. (2019). Lenia: Biology of artificial life. *Complex Systems*, 28(3), 251-286.
  6. Taleb, N. N. (2012). *Antifragile: Things that gain from disorder*. Random House.
  7. Holland, J. H. (1992). *Adaptation in natural and artificial systems*. MIT Press.
  8. Bedau, M. A. (2003). Artificial life: Organization, adaptation and complexity from the bottom up. *Trends in Cognitive Sciences*, 7(11), 505-512.
  9. Langton, C. G. (1989). *Artificial Life*. Addison-Wesley.
  10. Auer, P., Cesa-Bianchi, N., & Fischer, P. (2002). Finite-time analysis of the multiarmed bandit problem. *Machine Learning*, 47(2), 235-256.
  11. Moran, P. A. (1950). Notes on continuous stochastic phenomena. *Biometrika*, 37(1/2), 17-23.
  12. Shannon, C. E. (1948). A mathematical theory of communication. *Bell System Technical Journal*, 27(3), 379-423.
  13. Park, J. S., O'Brien, J. C., Cai, C. J., et al. (2023). Generative agents: Interactive simulacra of human behavior. *arXiv preprint arXiv:2304.03442*.
  14. Mnih, V., Kavukcuoglu, K., Silver, D., et al. (2015). Human-level control through deep reinforcement learning. *Nature*, 518(7540), 529-533.
  15. Gini, C. (1912). Variabilità e mutabilità. *Reprinted in Memorie di metodologica statistica*.
- 

## Appendix

---

### A. Full Experimental Results

[signoff.csv data for all seeds]

### B. Configuration Files

[YAML snippets]

### C. Reproducibility Protocol

For reviewers who wish to verify our results:

**1. Run Manifest** (`run-manifest.json`): Each experiment includes metadata for configuration verification: - Version identifier and configuration hash - Seed values and RNG isolation scheme - Timestamp and system information

**2. Verification Approach:** - **Statistical consistency:** Same seed + config produces equivalent statistical distributions - **Event sequence consistency:** Key emergence events (phase transitions, combo discoveries) occur at consistent time points - **Deterministic RNG isolation:** Separate generators for actions, spawning, and environment

**3. Challenge Seed Protocol:** 1. Run with published Challenge Seed (e.g., Seed 42) 2. Compare key metrics in `signoff.csv` (alive, combos, overall) 3. Verify event sequence in `events.jsonl` matches published milestones 4. Check lineage statistics in `blueprint_kpis.csv`

**4. Anti-Fabrication Guarantees:** - All metrics computed during simulation, not post-hoc - Event timestamps are monotonically increasing - Lineage trees are topologically consistent (no orphan nodes)

## D. Data Availability Statement

All experimental data supporting this paper is available in the supplementary materials package:

**Package Structure (~97MB total):**

```
Ordis_V3539_Data/
├── README.txt                                # Complete documentation
├── config_v3525_balanced.yaml                # Baseline condition config
├── config_v3529_winter.yaml                  # Winter condition config
|
├── V3533_signoff_winter_20x5k.csv          # Main results: 76% pass rate
├── V3533_signoff_baseline_20x5k.csv        # Main results: 6% pass rate
├── V3539_signoff_winter_20x5k.csv          # Extended: 58% pass rate
├── V3539_signoff_baseline_20x5k.csv        # Extended: 35% pass rate
|
├── V3533_winter_20x5k_raw/                  # Complete raw data (17 seeds)
│   └── seed_XXXX/
│       ├── events.jsonl                      # Per-seed directories
│       ├── tick_agg.jsonl                    # Event log (~1.2MB)
│       ├── sample_entities.jsonl            # Tick aggregation
│       ├── run-manifest.json               # Entity samples
│       └── blueprint_kpis.csv             # Run configuration
|
├── V3533_baseline_20x5k_raw/                # KPI metrics
|
└── historical_validation/
    ├── V347_ablation_study/              # Recovered validation experiments
    │   ├── L0_world_field_validation.json # 2×2×2 ANOVA (40 seeds)
    │   ├── L1_oims_validation.json       # 5/5 PASS
    │   ├── L2_genealogy_validation.json  # 5/5 PASS
    │   ├── L3_tds_dvs_validation.json   # 20/20 PASS
    │   └── H1_validation/              # 5/5 PASS
    |
    └── H1_validation/                  # Meta-rules validation
```

```
|── H234_validation/          # Communication/social validation  
└── combo_experiment/        # Combo system validation
```

**Data Availability:** Supplementary materials will be uploaded to Zenodo upon acceptance. DOI: [pending]

**License:** CC BY-NC 4.0 (NonCommercial)

---

**Document Type:** Technical Paper Skeleton (IMRaD) **Status:** Ready for expansion **Companion:** AGI海洛因\_完整报告.md (Chinese white paper version)